

script_PA

2025-08-11

Import the libraries.

```
library(tidyverse)

## — Attaching core tidyverse packages — tidyverse
## 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats    1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2     3.5.2      ✓ tibble     3.3.0
## ✓ lubridate  1.9.4      ✓ tidyr      1.3.1
## ✓ purrr      1.1.0
## — Conflicts —
tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all
## conflicts to become errors

library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method      from
## as.zoo.data.frame zoo

library(lubridate)
library(ggplot2)
library(scales)

##
## Attaching package: 'scales'
##
## The following object is masked from 'package:purrr':
##
##   discard
##
## The following object is masked from 'package:readr':
##
##   col_factor

library(readr)
library(dplyr)
```

Load the dataset.


```
data <- read_csv('CO2emissions.csv')
```

```
## Rows: 1488 Columns: 4
## — Column specification
##
## Delimiter: ","
## chr (2): Country, Code
## dbl (2): Year, Annual CO2 emissions (tonnes)
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.
```

General overview graph without forecasts.

```
data |>
  filter(!Country %in% c("World", "European Union (27)")) |>
  ggplot(aes(x = Year, y = `Annual CO2 emissions (tonnes)`, color = Country))
+
  geom_line(size = 1) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  labs(
    title = "Carbon Dioxide Emissions (tonnes) per Country",
    x = "Year",
    y = "",
    color = "Country"
  ) +
  theme_minimal()

## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```



Now we build a function that splits the time series into train (until 2019) and test (2020-2023) and compares the performances of ETS and ARIMA for every country. We evaluate the models using RMSE (Root Mean Squared Error).

```
compare_models <- function(df, country_name) {
  country_data <- df %>%
    dplyr::filter(Country == country_name) %>%
    arrange(Year)

  # Make it a ts
  ts_data <- ts(country_data$`Annual CO2 emissions (tonnes)`, start =
min(country_data$Year), frequency = 1)

  # Train/test split
  train_ts <- window(ts_data, end = 2019)
  test_ts <- window(ts_data, start = 2020, end = 2023)
  h <- length(test_ts)

  # ETS
  ets_fit <- ets(train_ts)
  # ETS(ANN) returns flat predictions -> so we force ETS(AAN)
  if (grepl("ETS\\(A,N,N\\)", ets_fit$method)) {
    ets_fit <- ets(train_ts, model = "AAN")
  }
}
```

```

}
ets_fc <- forecast(ets_fit, h = h)
ets_rmse <- sqrt(mean((test_ts - ets_fc$mean)^2))
ets_label <- ets_fit$method

# ARIMA
arima_fit <- auto.arima(train_ts, stepwise = FALSE, approximation = FALSE)
arima_order <- arimaorder(arima_fit)

# If ARIMA(0,1,0), then force a drift to avoid flat forecasts
if (all(arima_order == c(0, 1, 0))) {
  arima_fit <- Arima(train_ts, order = c(0, 1, 0), include.drift = TRUE)
}

arima_fc <- forecast(arima_fit, h = h)
arima_rmse <- sqrt(mean((test_ts - arima_fc$mean)^2))
arima_order <- arimaorder(arima_fit)
has_drift <- "drift" %in% names(coef(arima_fit))

# Best model
if (ets_rmse < arima_rmse) {
  best_model <- ets_label
} else {
  suffix <- if (has_drift) ")" + drift" else ")"
  best_model <- paste0("ARIMA(", arima_order[1], ",", arima_order[2], ",",
arima_order[3], suffix)
}

# Return the table
tibble(
  Country = country_name,
  ETS_RMSE = ets_rmse,
  ARIMA_RMSE = arima_rmse,
  Best_Model = best_model
)
}

```

Apply the function to all the countries.

```

# Make a list of all the countries
country_list <- data %>% distinct(Country) %>% pull(Country)

# See results
results <- map_dfr(country_list, ~compare_models(data, .x))
print(results)

## # A tibble: 12 × 4
##   Country          ETS_RMSE  ARIMA_RMSE Best_Model
##   <chr>          <dbl>      <dbl> <chr>
## 1 Austria      6471551.    6106282. ARIMA(0,1,4)

```

##	2	Belgium	13397039.	11184008.	ARIMA(1,1,3)
##	3	Denmark	1640563.	3078411.	ETS(M,A,N)
##	4	European Union (27)	165355401.	191969299.	ETS(A,A,N)
##	5	France	25176203.	25347771.	ETS(A,Ad,N)
##	6	Germany	72036682.	77315014.	ETS(A,Ad,N)
##	7	Italy	18389382.	18386440.	ARIMA(1,1,1)
##	8	Netherlands	25178780.	26491215.	ETS(A,A,N)
##	9	Poland	26505648.	21423190.	ARIMA(0,1,0) + drift
##	10	Spain	35705898.	22040784.	ARIMA(3,1,0)
##	11	Sweden	3018231.	4260719.	ETS(A,Ad,N)
##	12	World	1627139527.	1593542557.	ARIMA(0,2,2)

Now let's compute forecasts (2024-2030) for every country using the respective best model.

```

forecast_2030 <- function(df, country_name, best_model) {
  country_data <- df %>%
    dplyr::filter(Country == country_name) %>%
    arrange(Year)

  ts_data <- ts(country_data$`Annual CO2 emissions (tonnes)`, start =
min(country_data$Year), frequency = 1)

  h <- 2030 - 2023 # Years to forecast

  # ETS
  if (startsWith(best_model, "ETS")) {
    fit <- ets(ts_data)

    # Avoid ETS(A,N,N), brings to flat forecasts. Force ETS(A,A,N) instead
    if (grepl("ETS\\(A,N,N\\)", fit$method)) {
      fit <- ets(ts_data, model = "AAN")
    }

    fc <- forecast(fit, h = h)
    model_label <- fit$method
  }

  # ARIMA (with drift if (0,1,0))
  else if (startsWith(best_model, "ARIMA")) {
    drift <- grepl("\\+ drift$", best_model)

    # Extract parameters
    order_vals <- gsub("ARIMA\\(|\\)|\\+ drift", "", best_model)
    order_vec <- as.integer(strsplit(order_vals, ",")[[1]])

    fit <- Arima(ts_data, order = order_vec, include.drift = drift)
    fc <- forecast(fit, h = h)
  }
}

```

```

    model_label <- best_model
  }

  # Extract forecasts
  years <- as.character(2024:2030)
  preds <- as.numeric(fc$mean)
  tibble_row <- tibble(
    Country = country_name,
    Best_Model = model_label,
    !!!set_names(as.list(preds), years)
  )

  return(tibble_row)
}

```

Let's do it for all the countries.

```

forecast_table <- pmap_dfr(
  list(
    country_name = results$Country,
    best_model = results$Best_Model
  ),
  ~forecast_2030(data, ..1, ..2)
)

# Show table with best models
print(forecast_table)

## # A tibble: 12 × 9
##   Country      Best_Model `2024`  `2025`  `2026`  `2027`  `2028`  `2029`
##   <chr>         <chr>         <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
##   <dbl>
## 1 Austria      ARIMA(0,1,... 5.91e 7 5.72e 7 5.82e 7 5.90e 7 5.90e 7 5.90e 7
## 2 Belgium      ARIMA(1,1,... 8.23e 7 8.47e 7 8.33e 7 8.46e 7 8.34e 7 8.45e 7
## 3 Denmark      ETS(M,A,N)    2.64e 7 2.54e 7 2.44e 7 2.35e 7 2.25e 7 2.16e 7
## 4 European ... ETS(A,A,N)    2.45e 9 2.39e 9 2.33e 9 2.27e 9 2.21e 9 2.15e 9
## 5 France        ETS(A,Ad,... 2.69e 8 2.65e 8 2.62e 8 2.60e 8 2.58e 8 2.57e 8
## 6 Germany       ETS(A,Ad,... 5.98e 8 5.99e 8 5.99e 8 5.99e 8 6.00e 8 6.00e 8
## 7 Italy          ARIMA(1,1,... 3.08e 8 3.03e 8 2.99e 8 2.95e 8 2.91e 8 2.87e 8
## 8 Netherlan... ETS(A,A,N)    1.16e 8 1.14e 8 1.11e 8 1.08e 8 1.06e 8 1.03e 8
## 9 Poland        ARIMA(0,1,... 2.91e 8 2.93e 8 2.95e 8 2.97e 8 2.98e 8 3.00e 8

```

```

3.02e 8
## 10 Spain      ARIMA(3,1,... 2.22e 8 2.23e 8 2.20e 8 2.20e 8 2.20e 8 2.19e 8
2.19e 8
## 11 Sweden     ETS(A,Ad,... 3.63e 7 3.60e 7 3.57e 7 3.55e 7 3.53e 7 3.52e 7
3.50e 7
## 12 World      ARIMA(0,2,... 3.82e10 3.86e10 3.91e10 3.95e10 3.99e10 4.03e10
4.07e10

```

Now, build a general line chart with forecasts.

```


historical_df <- data %>%
  filter(Year >= 1990, Year <= 2023,
         !Country %in% c("World", "European Union (27)")) %>%
  select(Country, Year, Emissions = `Annual CO2 emissions (tonnes)`) %>%
  mutate(Forecast = FALSE)

# Forecasts for countries
forecast_long <- forecast_table %>%
  filter(!Country %in% c("World", "European Union (27)")) %>%
  pivot_longer(cols = `2024`:`2030`, names_to = "Year", values_to =
"Emissions") %>%
  mutate(Year = as.integer(Year),
         Forecast = TRUE)

# Join observed data and forecasts (not at a country-level)
full_df <- bind_rows(historical_df, forecast_long) %>%
  mutate(Forecast = ifelse(Forecast, "Forecasted", "Observed"))

# Plot the Linechart
ggplot(full_df, aes(x = Year, y = Emissions, color = Country, linetype =
Forecast)) +
  geom_line(size = 1) +
  scale_linetype_manual(values = c("Observed" = "solid", "Forecasted" =
"dashed")) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  labs(title = "CO2 Emissions (1990 - 2030)",
       x = "Year",
       y = "Emissions (tonnes)",
       color = "Country",
       linetype = "Type") +
  theme_minimal()

```



Now let's see the forecast more in detail for every country. Also visualizing the Fit for 55 target.

```
# Fit for 55 target
get_target_value <- function(ts_data) {
  ts_data[which(time(ts_data) == 1990)] * 0.45
}

# A df for each country
generate_country_forecast_df <- function(df, country_name, best_model) {
  country_data <- df %>%
    filter(Country == country_name) %>%
    arrange(Year)

  ts_data <- ts(country_data$`Annual CO2 emissions (tonnes)`, start =
min(country_data$Year), frequency = 1)
  h <- 2030 - 2023

  # Target
  fit_for_55 <- get_target_value(ts_data)

  # Re-Fit models
  if (startsWith(best_model, "ETS")) {
```



```

fit <- ets(ts_data)
if (grepl("ETS\\(A,N,N\\)", fit$method)) {
  fit <- ets(ts_data, model = "AAN")
}
fc <- forecast(fit, h = h, level = 95)
} else if (startsWith(best_model, "ARIMA")) {
  drift <- grepl("\\*drift$", best_model)
  order_vals <- gsub("ARIMA\\(|\\)|\\+ drift", "", best_model)
  order_vec <- as.integer(strsplit(order_vals, ",")[[1]])
  fit <- Arima(ts_data, order = order_vec, include.drift = drift)
  fc <- forecast(fit, h = h, level = 95)
}

# Observed data
observed_df <- tibble(
  Year = as.integer(time(ts_data)),
  Emissions = as.numeric(ts_data),
  Forecast = FALSE,
  Lower = NA,
  Upper = NA,
  Country = country_name,
  Fit55 = fit_for_55
)

# Forecasted data with 95% intervals
forecast_df <- tibble(
  Year = 2024:2030,
  Emissions = as.numeric(fc$mean),
  Lower = as.numeric(fc$lower[, 1]),
  Upper = as.numeric(fc$upper[, 1]),
  Forecast = TRUE,
  Country = country_name,
  Fit55 = fit_for_55
)

# Join observed and forecasted values for each country
bind_rows(observed_df, forecast_df)
}

forecast_plot_data <- pmap_dfr(
  list(
    country_name = results$Country,
    best_model = results$Best_Model
  ),
  ~generate_country_forecast_df(data, ..1, ..2)
)

plot_country_forecast <- function(country_name, save = FALSE) {
  df_country <- forecast_plot_data %>% filter(Country == country_name)

```

```

p <- ggplot(df_country, aes(x = Year, y = Emissions)) +
  geom_ribbon(
    data = df_country %>% filter(Forecast == TRUE),
    aes(x = Year, ymin = Lower, ymax = Upper),
    fill = "lightblue", alpha = 0.2
  ) +
  geom_line(aes(color = Forecast, linetype = Forecast), size = 1) +
  geom_hline(aes(yintercept = Fit55, linetype = "Fit for 55"), color = "red",
    linewidth = 0.8) +
  scale_color_manual(values = c("FALSE" = "black", "TRUE" = "blue"),
    labels = c("Observed", "Forecast")) +
  scale_linetype_manual(values = c("FALSE" = "solid", "TRUE" = "dashed",
    "Fit for 55" = "dashed"), breaks = c("Observed", "Forecast", "Fit for 55")) +
  scale_x_continuous(limits = c(1990, 2030)) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale()))
+
  labs(
    title = paste0("Forecast for ", country_name, " (95% Confidence)"),
    x = "Year",
    y = "Emissions (tonnes)",
    color = "",
    linetype = ""
  ) +
  theme_classic() +
  theme(
    panel.border = element_rect(color = "black", fill = NA, linewidth = 1),
    legend.position = "top",
    legend.justification = "center",
    legend.direction = "horizontal",
    legend.title = element_blank(),
    legend.text = element_text(size = 9),
    legend.spacing.x = unit(5, "pt"),
    plot.title = element_text(hjust = 0.5, size = 14, face = "bold"),
    plot.title.position = "plot",
    plot.margin = margin(t = 20, r = 10, b = 10, l = 10)
  )


return(p)
}

# Print chart for every country
for (country in country_list) {
  print(plot_country_forecast(country))
}

## Warning: Removed 90 rows containing missing values or values outside the
scale range
## (`geom_line()`).


```

Forecast for Austria (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Belgium (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Denmark (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for European Union (27) (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for France (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Germany (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```


Forecast for Italy (95% Confidence)



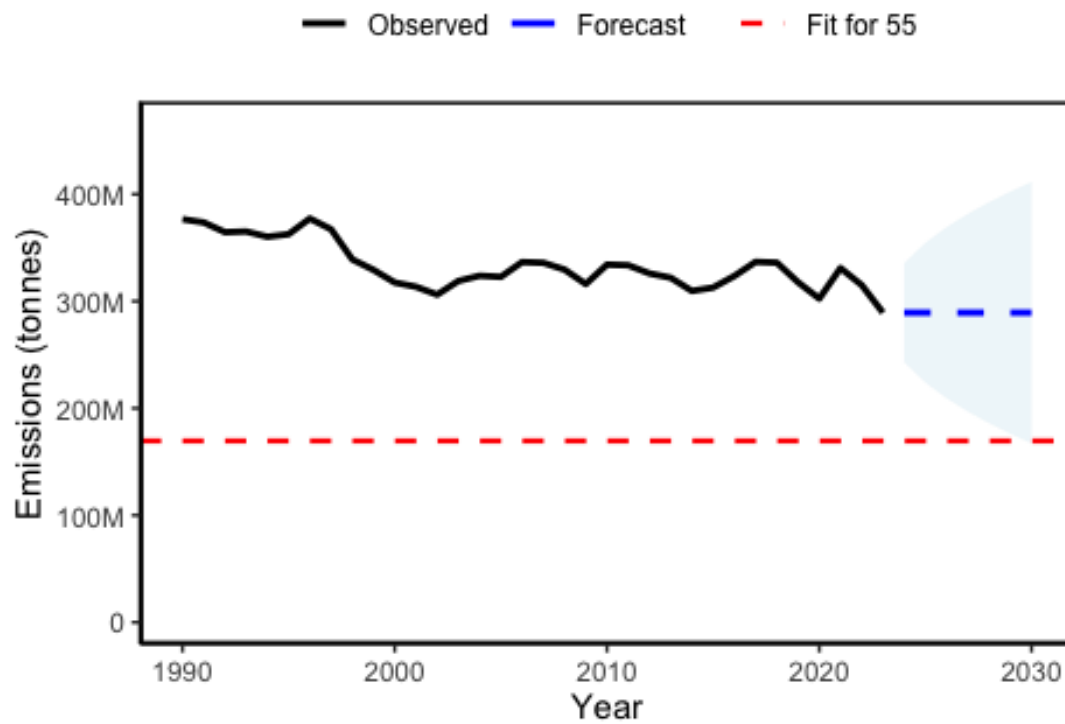
```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Netherlands (95% Confidence)




```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Poland (95% Confidence)



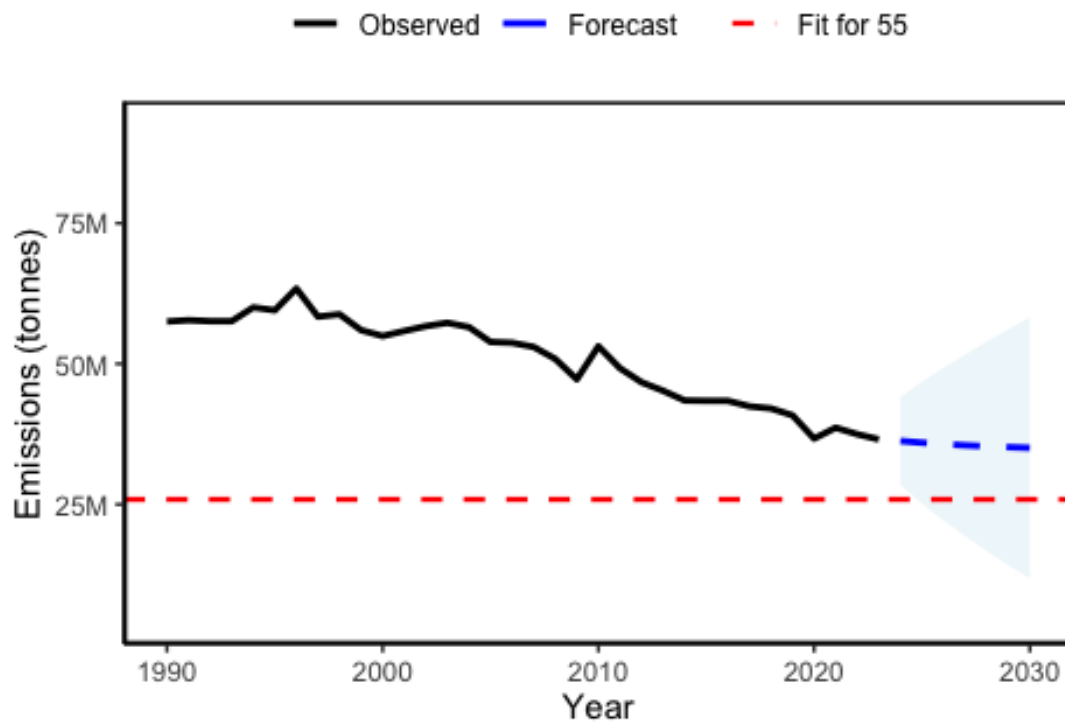
```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Spain (95% Confidence)



```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for Sweden (95% Confidence)



```
## Warning: Removed 90 rows containing missing values or values outside the  
scale range  
## (`geom_line()`).
```

Forecast for World (95% Confidence)

